

# **Agentic Programming -**

## **A Compiler Case Study**

### **Driving Effective Agentic Programming**

Jian Weng

CEMSE, KAUST

Week-1 Session-2

# I forget one thing last time

- But I did not want to do it last time, as some of you may not officially enroll the class yet...
- I also found it is super hard to interact with you without knowing you...
- Self-introduction around all of us:
  - Name
  - Background
  - Experience with AI coding tools
  - Expectation from this class

# Recap from Last Session

- This is NOT a class for vibe coding
- AI is giving you a **team of you**
  - AI is only as good as you are
  - But you save the most bandwidth on communication
- Two key pains of agentic programming:
  - i. Deep **human involvement** still needed
  - ii. How to manage a large codebase **effectively and reliably?**

# Today's Focus

- Pain #1 (human involvement) — we will address this later
- To reduce human involvement, we first need to solve:
  - **How to make AI write high-quality code?**
- Once the code quality is reliable, then we can reduce human oversight

# Agenda

- The problem with AI-generated code
- SDD: Spec/Standard-Driven Development
- DDD: Document-Driven Development
- TDD: Test-Driven Development
- The complete workflow
- Wrap-up

# The Problem: Before AI Agents

Each feature request used to take you a whole day:

- **Implement** – half the time
- **Test** – half the time (when buggy); zero (when it works)
- **Document** – "Why do I even do this? I remember everything..."

The goal: **move on faster**

But moving on faster ≠ skipping quality

# The Problem: With AI Agents

AI writes code for you, but:

- Maybe the code is **buggy**?
  - → Test it! Tests don't make it 100% correct, but better than nothing
- You didn't write it, you didn't even read it! How do you **understand** it?
  - → Documentation!
  - To my observations: documentation significantly helps AI understand
- Without tests and docs, you're flying blind

# Proposed Solution: Three Pillars

- **SDD** – Spec/Standard-Driven Development
- **DDD** – Document-Driven Development
- **TDD** – Test-Driven Development

The standard: **Write docs first, write tests second, write code last**

# Agenda

- The problem with AI-generated code
- **SDD: Spec/Standard-Driven Development**
- DDD: Document-Driven Development
- TDD: Test-Driven Development
- The complete workflow
- Wrap-up

# SDD: The Spec is the Charter

A **spec** defines how development must proceed:

- Coding standard
- Documentation standard
- Testing standard

The spec enforces:

1. AI reads the **spec of the development flow**
2. AI reads the **spec of the feature request**
3. AI makes a **plan**
4. AI **executes** the plan following the spec

# SDD: What Does the Standard Say?

The development standard we adopt:

1. **Document first** — update docs before writing any code
2. **Test second** — write test cases based on the documented interface
3. **Code last** — implement the code to pass the tests

This is the order. Not the other way around.

# SDD: Where Does the Spec Live?

Remember `CLAUDE.md` from last session?

`CLAUDE.md` / `AGENTS.md` — only the basics:

- How to **build** the project
- How to **run tests**
- Project overview (what this repo is)

**Planning & execution standards** live elsewhere:

- `commands/` — encode the development workflow as slash commands
- `rules/` — enforce coding, doc, and testing standards per file

`CLAUDE.md` = context. Commands & Rules = standards.

# SDD: Not Just Claude Code

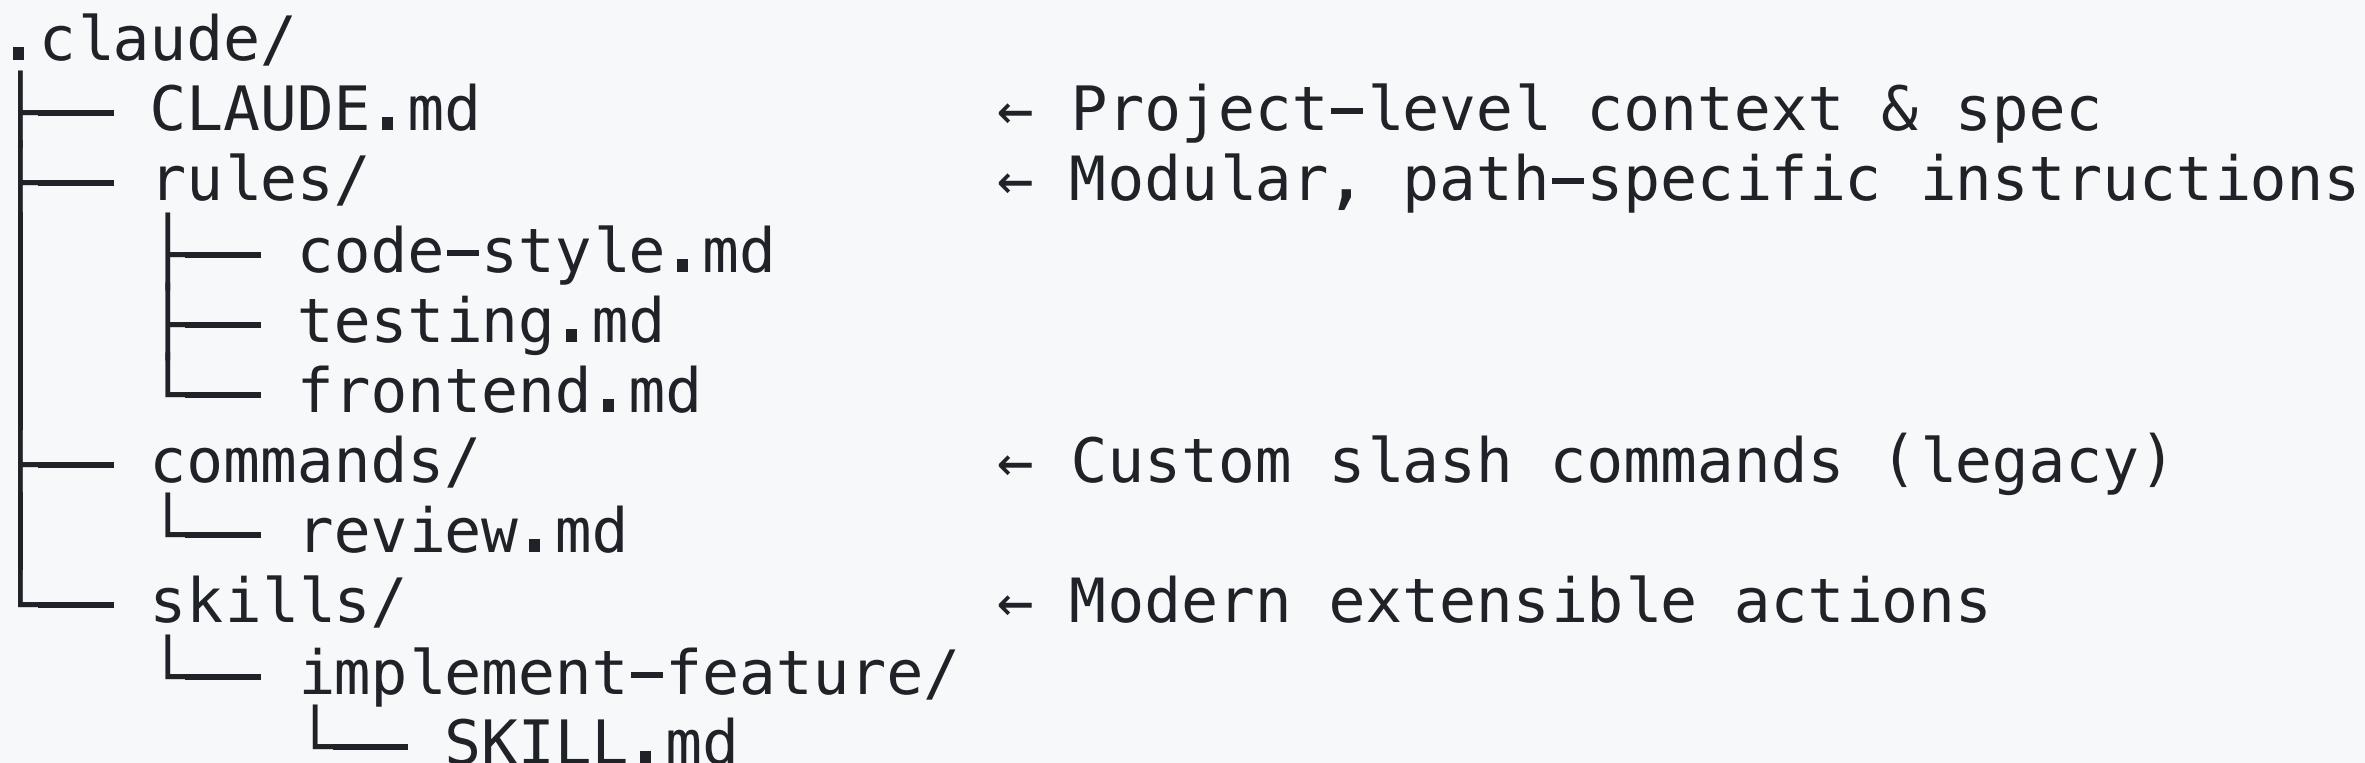
This layered system is **universal** across AI coding agents:

Claude Code	Codex	Cursor
CLAUDE.md	AGENTS.md	AGENTS.md
.claude/commands/	—	—
.claude/rules/	—	.cursor/rules/

- The names differ, but the **concept is the same**
- We use Claude Code as our example — it was the first AI coding CLI
  - It defined most of these standards
  - Other tools adopted similar patterns

# SDD: Beyond CLAUDE.md — Rules, Commands, Skills

Claude Code provides a layered system for encoding standards:



# SDD: Anatomy of Rules, Commands, and Skills

All of them share the same two-part structure:

---

- description: When does this apply?
- argument-hint: What arguments does it take?
- tools: What tools can it use?

---

(prompt: natural language instructions)

- What should the AI actually do?
- Step-by-step workflow

**Front matter** = machine-readable config (scope, triggers, options)

**Prompt** = human-readable instructions (the actual standard)

# SDD: Rules – Path-Specific Standards

Rules are markdown files in `.claude/rules/` that apply **conditionally**:

```
# .claude/rules/testing.md
```

```
---
```

```
paths:
```

- "src/\*\*/\*.cc"
- "src/\*\*/\*.h"

```
---
```

```
# Testing Rules
```

- Every public function must have a corresponding test
- Test files are named <module>\_test.cc
- Use GoogleTest framework
- Run tests with: make test

Rules with `paths` only activate when AI touches matching files.

# SDD: Commands — Reusable Workflows

Commands are single-file slash commands in `.claude/commands/`:

```
# .claude/commands/implementation.md
```

Read the feature request in \$ARGUMENTS.

Follow this workflow:

1. Read docs/ to understand the current design
2. Update the relevant documentation first
3. Write test stubs that compile but fail
4. Implement the code to pass all tests
5. Run `make test` to verify

Usage: `/implement "add string literal support to lexer"`

Commands encode your **development workflow** as a reusable action

# SDD: Skills – The Modern Standard

Skills are directories in `.claude/skills/` with richer capabilities:

```
# .claude/skills/document-module/SKILL.md
```

```
---
```

```
name: document-module
```

```
description: Generate module documentation
```

```
argument-hint: "[module path]"
```

```
---
```

For the module at \$ARGUMENTS, generate a `README.md` with:

1. **\*\*Purpose\*\*** – what this module does
2. **\*\*Public Interface\*\*** – exported functions and classes
3. **\*\*Internal Helpers\*\*** – private functions
4. **\*\*Data Structures\*\*** – types, enums, structs

Usage: Should be automatic when determined to be needed.

# SDD: Commands vs Skills

Think of them in C terms:

- **Command = a function call**
  - You invoke it explicitly: `/implement "add string literals"`
  - You pass arguments in, it runs a workflow
  - It is an **entry point** you control
- **Skill = a macro**
  - It expands automatically when the AI determines it's needed
  - Skills **cannot call commands** — by design
  - This keeps skills framework-agnostic and composable

# SDD: The Context Cost

Everything has a cost — **descriptions are loaded at session start:**

- Every `CLAUDE.md`, every rule, every skill `description`  
→ fed into the context **before you even type**
- This is your **startup overhead**

Implications:

- Write descriptions **concisely** — every token counts
- Too many verbose rules = bloated context = less room for actual work
- Think of it as: your spec is **always in memory**

Good specs are short and precise. Bad specs waste your context window.

# Anti-pattern: Vague Spec

What happens when the spec is too vague?

## # Bad CLAUDE.md

- Write good code
- Add tests when needed
- Document important things

AI interprets "good", "when needed", "important" **differently every time**

Result: inconsistent code, missing tests, sparse docs

| Be specific. The spec is a **contract**, not a suggestion.

# Agenda

- The problem with AI-generated code
- SDD: Spec/Standard-Driven Development
- **DDD: Document-Driven Development**
- TDD: Test-Driven Development
- The complete workflow
- Wrap-up

# DDD: Document as a Source Tree

In the AI era, documentation is not an afterthought.

**Two levels of documentation:**

- `docs/` — **architecture-level**: project overview, module relationships, design decisions
- Next to each source file — **file-level**: interfaces, helpers, data structures

Each `.cc` / `.h` file has a companion `.md` **right beside it**.

# DDD: Example – Doc Source Tree

```
project/
└── docs/
    └── architecture.md      ← Project architecture & module relationships
└── src/
    ├── lexer/
    │   ├── lexer.h
    │   ├── lexer.cc
    │   └── lexer.md          ← Public API, helpers, data structures
    ├── parser/
    │   ├── parser.h
    │   ├── parser.cc
    │   └── parser.md          ← Public API, helpers, data structures
    └── codegen/
        ├── codegen.h
        ├── codegen.cc
        └── codegen.md          ← Public API, helpers, data structures
```

docs/ = the big picture. src/\*.md = the details, right next to the code.

# DDD: What Goes in a Module Doc?

## # Lexer Module

### ## Purpose

Tokenizes source code into a stream of tokens.

### ## Public Interface

- `Lexer(std::string source)` – constructor
- `Token nextToken()` – returns next token, advances cursor
- `std::vector<Token> tokenize()` – tokenizes entire source

### ## Internal Helpers

- `skipWhitespace()` – advances past whitespace
- `readIdentifier()` – reads an identifier token
- `readNumber()` – reads a numeric literal

### ## Data Structures

- `Token { TokenType type; std::string value; int line; }`
- `enum TokenType { IDENT, NUMBER, PLUS, ... }`

# DDD: Who Writes the Docs?

- AI can help you write them
- You can write them manually
- Either way, **docs must exist before code changes**

Why document first?

- Updating the doc = **designing the interface**
- Like C's separation of `.h` (header) and `.c` (implementation)
- Once the interface is documented, you can write stubs and tests

Docs save AI's chain-of-thought when understanding your code later

# DDD: The Workflow

When a new feature request comes in:

1. **Update the doc** — add/modify the interface description
2. Now you have a clear picture of:
  - What functions exist
  - What parameters they take
  - What they return
3. The doc becomes the **blueprint** for stubs and tests

This is not extra work — this IS the design phase

# Anti-pattern: Code First, Doc Later

What actually happens when you say "I'll document later":

1. Write code → ship it → move on
2. Next feature comes in → "I'll read the code to understand"
3. AI reads 2000 lines → gets confused → hallucinates
4. You debug AI's hallucination → more time wasted than documenting

**With AI agents, the cost of missing docs is amplified**

- You don't just slow down yourself
- You slow down every future AI interaction with that code

# Agenda

- The problem with AI-generated code
- SDD: Spec/Standard-Driven Development
- DDD: Document-Driven Development
- **TDD: Test-Driven Development**
- The complete workflow
- Wrap-up

# Traditional TDD vs Agent TDD

**Traditional TDD (human writes everything):**

1. Human writes a failing test
2. Human writes code to pass the test
3. Human refactors
4. Repeat

**Agent TDD (human designs, AI implements):**

1. Human designs the interface (via docs)
2. Human (or AI) writes stubs + failing tests
3. AI implements code to pass the tests

# Why Agent TDD is More Powerful

Traditional TDD:

- Human must context-switch between test-writing and implementation
- Temptation to skip tests when "the code obviously works"

Agent TDD:

- Clear separation of concerns
  - **You:** design interface, write tests (the spec of correctness)
  - **AI:** implement until tests pass
- Tests are **not optional** — they are the AI's success criterion
- No more "it obviously works" — prove it

# TDD: Step by Step

After docs define the interfaces:

1. Write **empty function stubs** that match the documented interface
2. Write **test cases** that call these stubs
3. Tests **compile** — but they **fail** (stubs return wrong/no values)
4. Now tell AI: **make the tests pass**

```
// stub – compiles, but tests will fail
Token Lexer::nextToken() {
    return Token{TokenType::UNKNOWN, "", 0}; // TODO
}
```

# TDD: Example – Test Cases

```
TEST(LexerTest, SingleNumber) {
    Lexer lexer("42");
    Token tok = lexer.nextToken();
    EXPECT_EQ(tok.type, TokenType::NUMBER);
    EXPECT_EQ(tok.value, "42");
}
```

```
TEST(LexerTest, SimplePlus) {
    Lexer lexer("1 + 2");
    auto tokens = lexer.tokenize();
    ASSERT_EQ(tokens.size(), 3);
    EXPECT_EQ(tokens[0].type, TokenType::NUMBER);
    EXPECT_EQ(tokens[1].type, TokenType::PLUS);
    EXPECT_EQ(tokens[2].type, TokenType::NUMBER);
}
```

# TDD: The Prompt to AI

Now you tell the agent:

Run `make test`. Some tests are failing.

Read the documentation in `docs/lexer/README.md` to understand the expected behavior.

Implement the code in `src/lexer/lexer.cc` to make all tests pass.

The agent has:

- A **clear goal** (tests pass)
- A **clear reference** (documentation)
- A **clear scope** (specific files)

# TDD: Leveraging Agent Persistence

A key insight about AI agents:

- All agents are **persistent** — they keep iterating toward a goal
- If the goal is not met, they will try again and again
- We **leverage** this property:
  - Give the agent a **clear, verifiable target** (all tests pass)
  - Let it grind until it succeeds

The agent won't complain. It won't get tired.

It will keep fixing bugs until the tests are green.

# Anti-pattern: No Tests, Just "Looks Right"

Without tests, the AI agent has no feedback loop:

1. AI writes code
2. You read it → "looks right" → ship it
3. Bug surfaces in production → back to square one

Or worse:

1. AI writes code
2. You don't read it → ship it
3. Bug surfaces → you don't even know where to look

Tests are the **minimum viable verification**.

# Anti-pattern: Writing Tests After Code

Why not write code first, then add tests?

- AI writes code → it "works" → you write tests that match the code
- **Circular reasoning:** tests verify what the code does, not what it should do
- You end up testing the implementation, not the specification

The correct order:

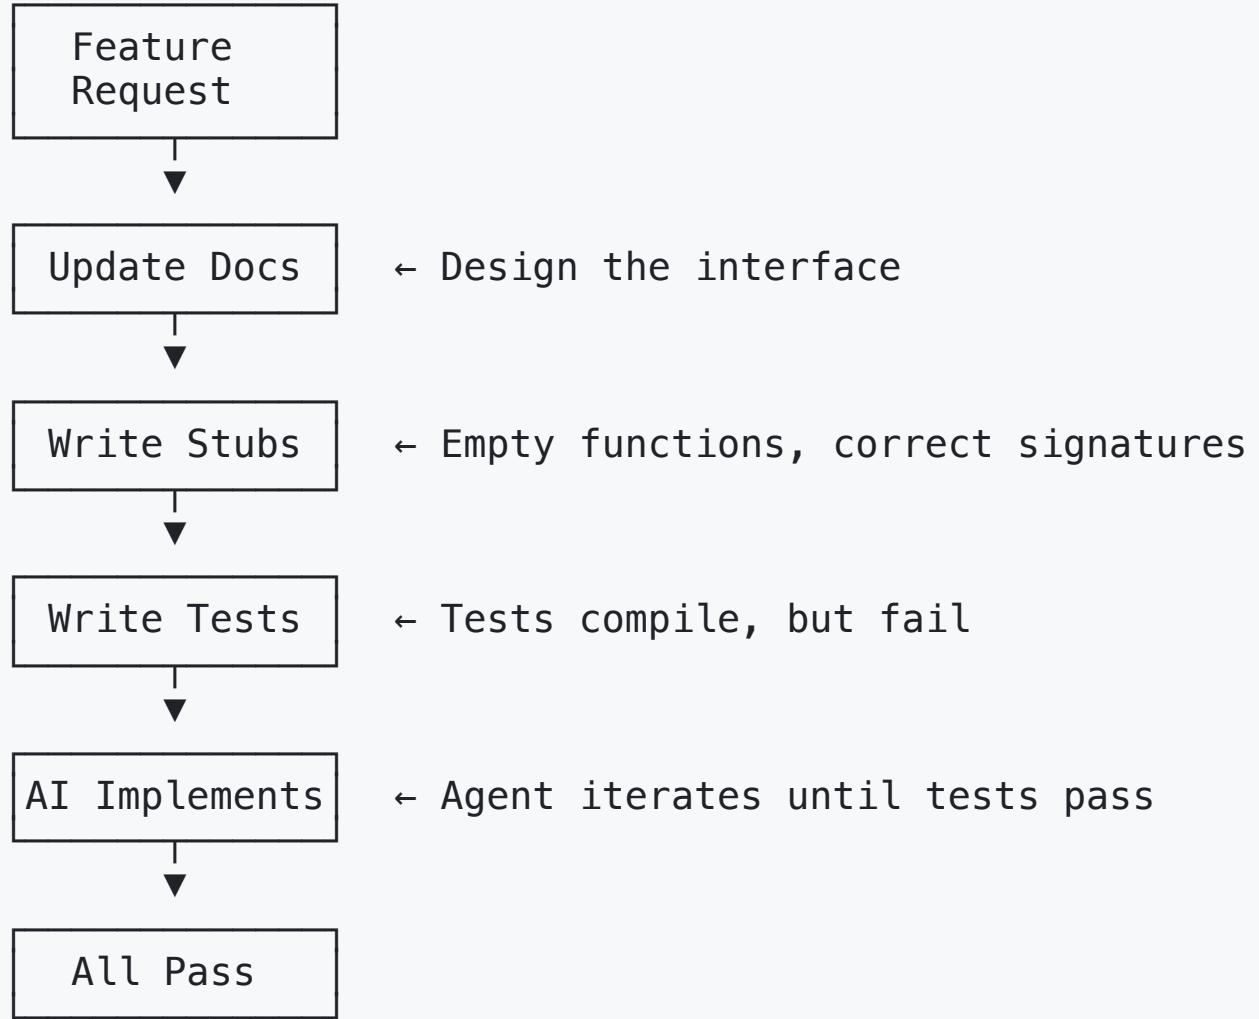
1. Spec defines **what should happen**
2. Tests encode **what should happen**
3. Code makes it happen

Tests written after code are just a rubber stamp.

# Agenda

- The problem with AI-generated code
- SDD: Spec/Standard-Driven Development
- DDD: Document-Driven Development
- TDD: Test-Driven Development
- **The complete workflow**
- Wrap-up

# The Complete Workflow



# Walkthrough: Adding a Feature

Let's say we want to add **string literal support** to our lexer.

## Step 1: Update docs

### **## Public Interface (updated)**

- `Token nextToken()` – now also handles string literals

### **## Data Structures (updated)**

- `enum TokenType { ..., STRING, ... }`

### **## Behavior**

- String literals are enclosed in double quotes: "hello"
- Escape sequences supported: \n, \t, \\, \"

# Walkthrough: Stubs and Tests

## Step 2: Write stub

```
// lexer.cc – string case added, returns UNKNOWN for now
case '':
    return Token{TokenType::UNKNOWN, "", line_};
```

## Step 3: Write tests

```
TEST(LexerTest, StringLiteral) {
    Lexer lexer("\"hello\"");
    Token tok = lexer.nextToken();
    EXPECT_EQ(tok.type, TokenType::STRING);
    EXPECT_EQ(tok.value, "hello");
}
```

```
TEST(LexerTest, StringEscape) {
```

# Walkthrough: AI Implements

## Step 4: Prompt the agent

The string literal tests in `lexer_test.cc` are failing.

Read `docs/lexer/README.md` for the expected behavior.

Implement string literal tokenization in `lexer.cc`.

Run `make test` to verify.

### What happens:

1. Agent reads the doc → understands escape sequences
2. Agent implements `readString()` helper
3. Runs tests → some fail (forgot `\\"` escape)
4. Fixes → runs again → all pass

# Why This Works

- **Docs** give the agent context without guessing
- **Stubs** define the exact interface to implement
- **Tests** provide a clear, automated success criterion
- **Agent persistence** means it will keep trying until it works
- **You** stay in control of design, but delegate implementation

# Putting It All Together: CLAUDE.md

Your `CLAUDE.md` ties everything together:

## # My Compiler Project

### ## Build & Test

- Build: make
- Test: make test

### ## Development Standard

1. Update docs in docs/ BEFORE writing code
2. Write test cases that compile but fail
3. Implement code to make tests pass

### ## Documentation Standard

- Each module has docs/<module>/README.md
- Document: purpose, public API, helpers, data structures

# Wrap-up

- High-quality AI code requires structure: **SDD + DDD + TDD**
- Document first → Test second → Code last
- Documentation mirrors your source tree
- Tests give agents a clear, verifiable goal
- Agents are persistent — leverage that property
- `CLAUDE.md` encodes both context and standards

## Next session

- We will put this into practice
- Setting up the spec and documentation for our compiler project